

## ASSIGNMENT -2

### Data Visualization and Pre-processing

Assignment Date	26 September 2022
Team ID	PNT2022TMID45335
Project Name	AI BASED DISCOURSE FOR BANKING INDUSTRY
Student Name	NATHANIEL NICHOLAS J S
Student Roll Number	E1195043
Maximum Marks	2 Marks

#### Question 1

- 1.Download the Data set
- 2.Load The Dataset

#### Solution:

```
import numpy as np
```

```
import pandas as pd
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
import sklearn
```

```
data = pd.read_csv(r'C:\Users\ADMIN\Downloads\Churn_Modelling.csv')
```

```
data.head()
```

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import sklearn
data = pd.read_csv(r'C:\Users\ADMIN\Downloads\Churn_Modelling.csv')
data.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

### Question-3

3.Perform Below Visualizations

#### Solution:


##### #1.Univariate Analysis

```
sns.histplot(data['CreditScore'])
```

```
#1.Univariate Analysis
```


```
sns.histplot(data['CreditScore'])
```

```
<AxesSubplot:xlabel='CreditScore', ylabel='Count'>
```



```
sns.boxplot(x=data['Age'])
```

```
<AxesSubplot:xlabel='Age'>
```




## #2.Bivariate Analysis

```
plt.figure(figsize=(15,15))  
sns.barplot(x=data['Age'],y=data['EstimatedSalary'])
```

#2.Bivariate Analysis

```
plt.figure(figsize=(15,15))  
sns.barplot(x=data['Age'],y=data['EstimatedSalary'])
```


<AxesSubplot:xlabel='Age', ylabel='estimatedSalary'>



plt.figure(figsize=(15,15))


```
plt.figure(figsize=(10,10))
sns.barplot(x=data['CreditScore'],y=data['Tenure'])
```

<AxesSubplot:xlabel='CreditScore', ylabel='Tenure'>




```
plt.figure(figsize=(10,10))
sns.lineplot(x=data['Balance'],y=data['Age'])
```


<AxesSubplot:xlabel='Balance', ylabel='Age'>



```
plt.figure(figsize=(10,10))
sns.lineplot(x=data['Age'],y=data['CreditScore'])
<AxesSubplot:xlabel='Age', ylabel='CreditScore'>
```



```
plt.figure(figsize=(10,10))
sns.scatterplot(x=data['Age'],y=data['CreditScore'])
<AxesSubplot:xlabel='Age', ylabel='CreditScore'>
```




### #3.Multi variate analysis

```
sns.pairplot(data= data [['Age','EstimatedSalary','Tenure']],hue='Age')
```

```
#3.Multi variate analysis
```


```
sns.pairplot(data= data [['Age','EstimatedSalary','Tenure']],hue='Age')
```

```
<seaborn.axisgrid.PairGrid at 0x2c7679deb50>
```



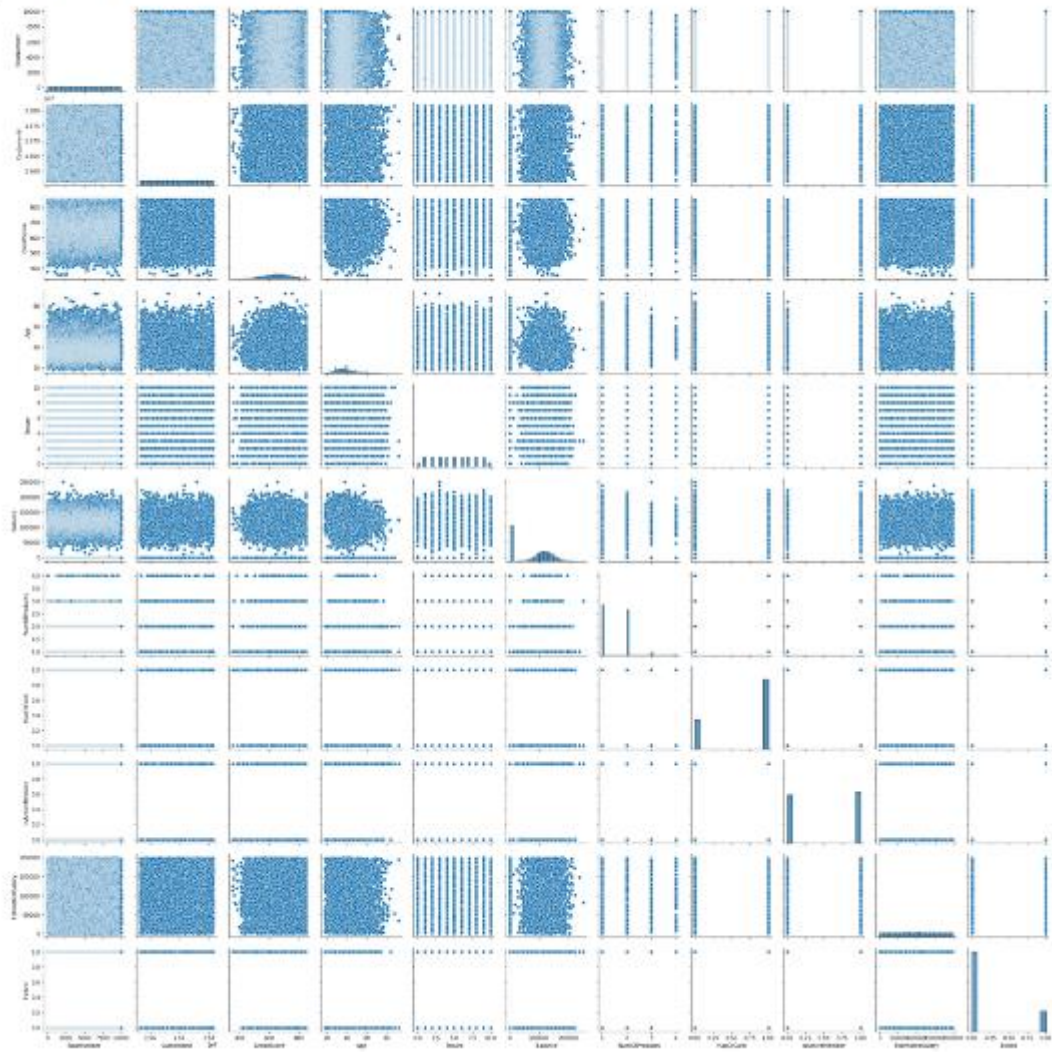
```
sns.scatterplot(x=data['Tenure'],y=data['CreditScore'], hue = data['Gender'])
```

```
<AxesSubplot:xlabel='Tenure', ylabel='CreditScore'>
```



```
IN [110]: sns.pairplot(data)
```

```
Out[110]: <seaborn.pairgrid.PairGrid at 0x1c788cc17f0>
```





#### Question.4:

Perform descriptive statistics on the dataset

#### Solution:

```
data.mean(numeric_only = True)
```

```
data.mean(numeric_only = True)
```

RowNumber	5.000500e+03
CustomerId	1.569094e+07
CreditScore	6.505288e+02
Age	3.892180e+01
Tenure	5.012800e+00
Balance	7.648589e+04
NumOfProducts	1.530200e+00
HasCrCard	7.055000e-01
IsActiveMember	5.151000e-01
EstimatedSalary	1.000902e+05
Exited	2.037000e-01

dtype: float64

```
data.median(numeric_only = True)
```

```
In [112]: data.median(numeric_only = True)
```

```
Out[112]:
```

RowNumber	5.000500e+03
CustomerId	1.569074e+07
CreditScore	6.520000e+02
Age	3.700000e+01
Tenure	5.000000e+00
Balance	9.719854e+04
NumOfProducts	1.000000e+00
HasCrCard	1.000000e+00
IsActiveMember	1.000000e+00
EstimatedSalary	1.001939e+05
Exited	0.000000e+00

dtype: float64

```
data['CreditScore'].mode()  
data['Age'].mode()  
data['Balance'].unique()  
data['Tenure'].unique()  
data.std(numeric_only=True)
```

```
data['CreditScore'].mode()
```

```
0    850  
Name: CreditScore, dtype: int64
```

```
data['Age'].mode()
```

```
0    37  
Name: Age, dtype: int64
```

```
data['Balance'].unique()
```

```
array([ 0.,    83807.86, 159660.8, ...,  57369.61,  75075.31,  
       130142.79])
```

```
data['Tenure'].unique()
```

```
array([ 2,  1,  8,  7,  4,  6,  3, 10,  5,  9,  0], dtype=int64)
```

data.describe()

```
data.describe()
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.00000	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.515100	100090.239881	0.203700
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.499797	57510.492818	0.402765
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000000	11.580000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.000000	51002.110000	0.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000	1.000000	100193.915000	0.000000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000	1.000000	149388.247500	0.000000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.00000	1.000000	199992.480000	1.000000

data['NumOfProducts'].value\_counts()

```
In [119]: data['NumOfProducts'].value_counts()
```

```
Out[119]: 1    5084  
          2    4590  
          3     266  
          4      60  
          Name: NumOfProducts, dtype: int64
```

### Question.5

Handle the Missing values

#### Solution:

```
data.isnull().any()
```

```
data.isnull().sum()
```

```
In [120]: data.isnull().any()
```

```
Out[120]:
```

RowNumber	False
CustomerId	False
Surname	False
CreditScore	False
Geography	False
Gender	False
Age	False
Tenure	False
Balance	False
NumOfProducts	False
HasCrCard	False
IsActiveMember	False
EstimatedSalary	False
Exited	False

dtype: bool

```
In [121]: data.isnull().sum()
```

```
Out[121]:
```

RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0
NumOfProducts	0
HasCrCard	0
IsActiveMember	0
EstimatedSalary	0
Exited	0

dtype: int64

### Question.6


Find the outliers and replace the outliers

#### Solution:

```
sns.boxplot(x=data['Age'])
```

```
] : sns.boxplot(x=data['Age'])
```

```
] : <AxesSubplot:xlabel='Age'>
```



```
fig, ax = plt.subplots(figsize = (5,3)) #Outlier detection - Scatter plot  
ax.scatter(data['Balance'], data['Exited'])
```

```
# x-axis label  
ax.set_xlabel('Balance')
```

```
# y-axis label  
ax.set_ylabel('Exited')  
plt.show()
```


```
sns.boxplot(x=data['Balance'])
```

```
fig, ax = plt.subplots(figsize = (5,3)) #Outlier detection - Scatter plot
ax.scatter(data['Balance'], data['Exited'])


# x-axis label
ax.set_xlabel('Balance')

# y-axis label
ax.set_ylabel('Exited')
plt.show()

sns.boxplot(x=data['Balance'])
```



```
<AxesSubplot:xlabel='Balance'>
```



```
from scipy import stats #Outlier detection - zscore
zscore = np.abs(stats.zscore(data['CreditScore']))
print(zscore)
print('No. of Outliers : ', np.shape(np.where(zscore>3)))
```

```
from scipy import stats #Outlier detection - zscore
zscore = np.abs(stats.zscore(data['CreditScore']))
print(zscore)
print('No. of Outliers : ', np.shape(np.where(zscore>3)))
```

```
0      0.326221
1      0.440036
2      1.536794
3      0.501521
4      2.063884
...
9995    1.246488
9996    1.391939
9997    0.604988
9998    1.256835
9999    1.463771
Name: CreditScore, Length: 10000, dtype: float64
No. of Outliers : (1, 8)
```

```
q = data.quantile([0.70,0.30])
```

```
q
```

```
q = data.quantile([0.70,0.30])
q
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0.7	7000.3	15740461.6	704.0	42.0	7.0	122029.87	2.0	1.0	1.0	139432.236	0.0
0.3	3000.7	15641363.9	598.7	33.0	3.0	0.00	1.0	1.0	0.0	60736.079	0.0

```
iqr = q.iloc[0] - q.iloc[1]
```

```
iqr
```

```
iqr = q.iloc[0] - q.iloc[1]
iqr
```

```
RowNumber      3999.600
CustomerId      99097.700
CreditScore     105.300
Age              9.000
Tenure           4.000
Balance        122029.870
NumOfProducts    1.000
HasCrCard        0.000
IsActiveMember    1.000
EstimatedSalary  78696.157
Exited           0.000
dtype: float64
```

```
u = q.iloc[0] + (1.5*iqr)
u
```

```
u = q.iloc[0] + (1.5*iqr)
u
```

RowNumber	1.299970e+04
CustomerId	1.588911e+07
CreditScore	8.619500e+02
Age	5.550000e+01
Tenure	1.300000e+01
Balance	3.050747e+05
NumOfProducts	3.500000e+00
HasCrCard	1.000000e+00
IsActiveMember	2.500000e+00
EstimatedSalary	2.574765e+05
Exited	0.000000e+00

dtype: float64

```
l = q.iloc[1] - (1.5*iqr)
l
```

```
l = q.iloc[1] - (1.5*iqr)
l
```

RowNumber	-2.998700e+03
CustomerId	1.549272e+07
CreditScore	4.407500e+02
Age	1.950000e+01
Tenure	-3.000000e+00
Balance	-1.830448e+05
NumOfProducts	-5.000000e-01
HasCrCard	1.000000e+00
IsActiveMember	-1.500000e+00
EstimatedSalary	-5.730816e+04
Exited	0.000000e+00

dtype: float64

```
Q1 = data['EstimatedSalary'].quantile(0.30) #Outlier detection - IQR
Q3 = data['EstimatedSalary'].quantile(0.70)
iqr = Q3 - Q1
print(iqr)
upper=Q3 + 1.5 * iqr
lower=Q1 - 1.5 * iqr
count = np.size(np.where(data['EstimatedSalary'] > upper))
count = count + np.size(np.where(data['EstimatedSalary'] < lower))
print('No. of outliers : ', count)
```

```

Q1 = data['EstimatedSalary'].quantile(0.30) #Outlier detection - IQR
Q3 = data['EstimatedSalary'].quantile(0.70)
iqr = Q3 - Q1
print(iqr)
upper=Q3 + 1.5 * iqr
lower=Q1 - 1.5 * iqr
count = np.size(np.where(data['EstimatedSalary'] > upper))
count = count + np.size(np.where(data['EstimatedSalary'] < lower))
print('No. of outliers : ', count)

```

```

78696.157
No. of outliers : 0


```

```

data['CreditScore'] = np.where(np.logical_or(data['CreditScore']>900, data['CreditScore']<383), 650, data['CreditScore'])
sns.boxplot(data['CreditScore'])

```

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid FutureWarning  
<matplotlib.axes.\_subplots.AxesSubplot at 0x7f47eb744d90>



```

upper = data.Age.mean() + (3 * data.Age.std()) #Outlier detection - 3 sigma
lower = data.Age.mean() - (3 * data.Age.std())
columns = data[ ( data['Age'] > upper ) | ( data['Age']<lower ) ]
print('Upper range : ', upper)
print('Lower range : ', lower)
print('No. of Outliers : ', len(columns))

```

```

Upper range : 70.38521935511383
Lower range : 7.458380644886169
No. of Outliers : 133

```

columns = ['EstimatedSalary', 'Balance', 'Tenure'] #After outlier removal

for i in columns:

```

Q1 = data[i].quantile(0.30)
Q3 = data[i].quantile(0.70)
iqr = Q3 - Q1
upper=Q3 + 1.5 * iqr
lower=Q1 - 1.5 * iqr
count = np.size(np.where(data[i] > upper))
count = count + np.size(np.where(data[i] < lower))

```



```
print('No. of outliers in ', i, ' : ', count)
```

```
columns = ['EstimatedSalary', 'Balance', 'Tenure'] #After outlier removal

for i in columns:
    Q1 = data[i].quantile(0.30)
    Q3 = data[i].quantile(0.70)
    iqr = Q3 - Q1
    upper=Q3 + 1.5 * iqr
    lower=Q1 - 1.5 * iqr
    count = np.size(np.where(data[i] >upper))
    count = count + np.size(np.where(data[i] <lower))
    print('No. of outliers in ', i, ' : ', count)
```

```
No. of outliers in EstimatedSalary : 0
No. of outliers in Balance : 0
No. of outliers in Tenure : 0
```

## Question:7

### Check for Categorical columns and perform encoding

#### Solution:

```
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
le = LabelEncoder()
oneh = OneHotEncoder()
data['Surname'] = le.fit_transform(data['Surname'])
data['Gender'] = le.fit_transform(data['Gender'])
data['Geography'] = le.fit_transform(data['Geography'])
data.head()
```

```
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
le = LabelEncoder()
oneh = OneHotEncoder()
data['Surname'] = le.fit_transform(data['Surname'])
data['Gender'] = le.fit_transform(data['Gender'])
data['Geography'] = le.fit_transform(data['Geography'])
data.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	1115	619	0	0	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	1177	608	2	0	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	2040	502	0	0	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	289	699	0	0	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	1822	850	2	0	43	2	125510.82	1	1	1	79084.10	0

## Question.8

Split the data into dependent and independent variables split the data in X and Y

**Solution:**

```
x = data.iloc[:, 0:13]
```

x # independent values (inputs)

```
x = data.iloc[:, 0:13]
x # independent values (inputs)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	
	0	1	15634602	1115	619	0	0	42	2	0.00	1	1	1	101348.88
	1	2	15647311	1177	608	2	0	41	1	83807.86	1	0	1	112542.58
	2	3	15619304	2040	502	0	0	42	8	159660.80	3	1	0	113931.57
	3	4	15701354	289	699	0	0	39	1	0.00	2	0	0	93826.63
	4	5	15737888	1822	850	2	0	43	2	125510.82	1	1	1	79084.10
	...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	1999	771	0	1	39	5	0.00	2	1	0	0	96270.64
9996	9997	15569892	1336	516	0	1	35	10	57369.61	1	1	1	1	101699.77
9997	9998	15584532	1570	709	0	0	36	7	0.00	1	0	1	1	42085.58
9998	9999	15682355	2345	772	1	1	42	3	75075.31	2	1	0	0	92888.52
9999	10000	15628319	2751	792	0	0	28	4	130142.79	1	1	0	0	38190.78

10000 rows × 13 columns

```
y = data['Exited']
```

y # dependent values (output)

```
y = data['Exited']
y # dependent values (output)
```

```
0      1
1      0
2      1
3      0
4      0
..
9995   0
9996   0
9997   1
9998   1
9999   0
Name: Exited, Length: 10000, dtype: int64
```

### Question:9

Scale the independent variables

**Solution:**

```
from sklearn.preprocessing import StandardScaler, MinMaxScaler
sc = StandardScaler()
x_scaled = sc.fit_transform(x)
x_scaled
```

```
from sklearn.preprocessing import StandardScaler, MinMaxScaler
sc = StandardScaler()
x_scaled = sc.fit_transform(x)
x_scaled
```

```
array([[ -1.73187761, -0.78321342, -0.46418322, ...,  0.64609167,
         0.97024255,  0.02188649],
       [ -1.7315312 , -0.60653412, -0.3909112 , ..., -1.54776799,
         0.97024255,  0.21653375],
       [ -1.73118479, -0.99588476,  0.62898807, ...,  0.64609167,
        -1.03067011,  0.2406869 ],
       ...,
       [  1.73118479, -1.47928179,  0.07353887, ..., -1.54776799,
         0.97024255, -1.00864308],
       [  1.7315312 , -0.11935577,  0.98943914, ...,  0.64609167,
        -1.03067011, -0.12523071],
       [  1.73187761, -0.87055909,  1.4692527 , ...,  0.64609167,
        -1.03067011, -1.07636976]])
```

### Question:10

Split x and y into Training and Testing

**Solution:**

```
import pandas as pd
data=pd.read_csv("/content/Churn_Modelling.csv")
```

```
data.describe()
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.00000	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.515100	100090.239881	0.203700
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.499797	57510.492818	0.402769
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000000	11.580000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.000000	51002.110000	0.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000	1.000000	100193.915000	0.000000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000	1.000000	149388.247500	0.000000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.00000	1.000000	199992.480000	1.000000

```
import numpy as np
```

```
x=np.array(data["CustomerId"]).reshape(-1,1)
```

```
x.shape
```

```
import numpy as np
```

```
x=np.array(data["CustomerId"]).reshape(-1,1)
```

```
x.shape
```

```
(10000, 1)
```



```
y=np.array(data["EstimatedSalary"])
```

```
y.shape
```

```
(10000,)
```

```
[ ] print(y)
```

```
[101348.88 112542.58 113931.57 ... 42085.58 92888.52 38190.78]
```

```
[ ] print(type(x))
```

```
<class 'numpy.ndarray'>
```

```
[ ] from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30)
```

```
[ ] x_train.shape
```

```
(7000, 1)
```

```
x_test
```

```
array([[15611365],  
       [15610379],  
       [15641690],  
       ...,  
       [15724876],  
       [15765952],  
       [15661330]])
```

```
x_test.shape
```

```
(3000, 1)
```

```
y_test
```

```
array([ 60905.51, 121124.53, 163714.92, ...,  33245.97, 188382.77,  
       116141.72])
```

```
y_test.shape
```

```
(3000,)
```

```
y.shape
```

```
(10000,)
```

```
print(y_train.shape)
```

```
(7000,)
```

```
print(y_test.shape)
```

```
(3000,)
```